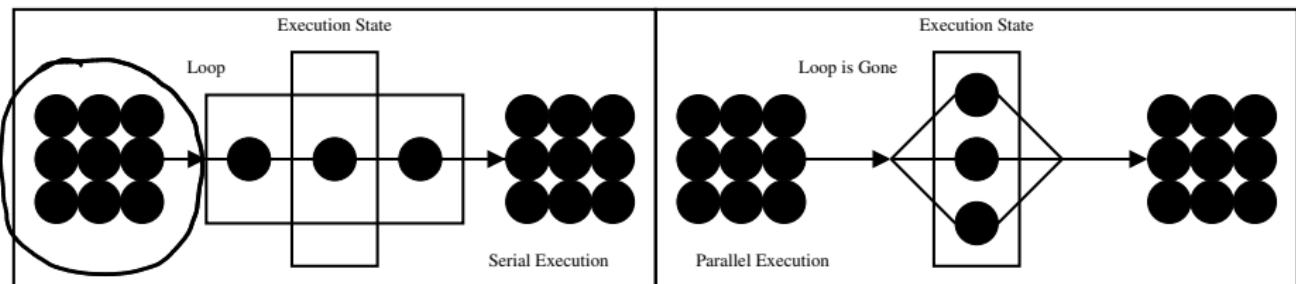


Serial vs Parallel vs Distributed (cont.)

Parallel Computing Systems

- Several processors that are located within a small distance of each other
- Their main purpose is to perform a computation task jointly
- Communication between processors, if required, is reliable, fast, and predictable.



Units of work

Figure 6: Difference between serial and parallel execution

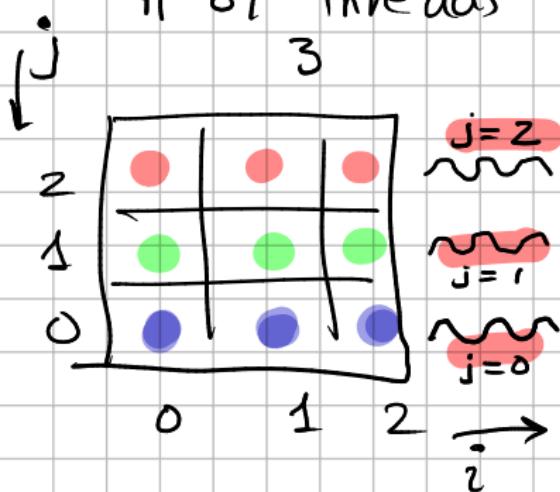
j → replace with thread.
 for ($j = 0$; $j < 3$; $j++$) $j = 0, 1, 2$

i for ($i = 0$; $i < 3$; $i++$) — serial

$a[i][j]$

3 threads I work / thread

$$\text{total work} = 3 \times 3 = 9$$

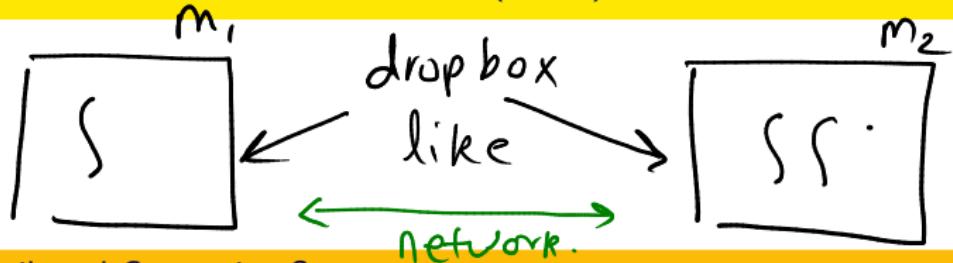


three threads start

$\text{for } (i = 0; i < 3; i++)$
 $a[i][j]$

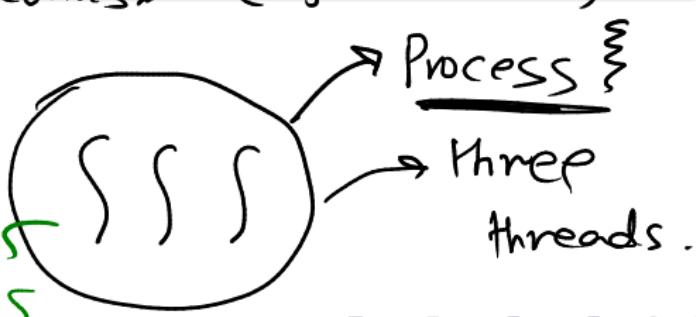
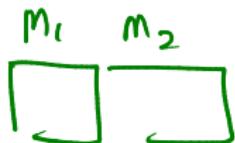
three threads end.

Serial vs Parallel vs Distributed (cont.)



- Processors may be far apart (are loosely coupled and usually independent of others)
- Challenges
 - Communication latency between processors is large and unpredictable
 - Communication links may be unreliable
 - Topology of system may undergo changes.
- Load Balancing difficulties → (synchronization)

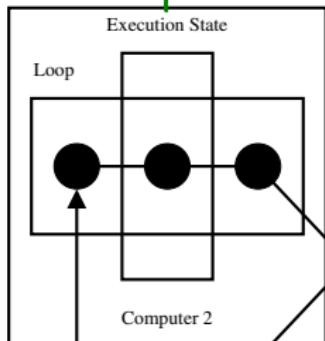
Work items > # of machines



Serial vs Parallel vs Distributed (cont.)

Sequential

1 core
machine



Parallel

3 threads.

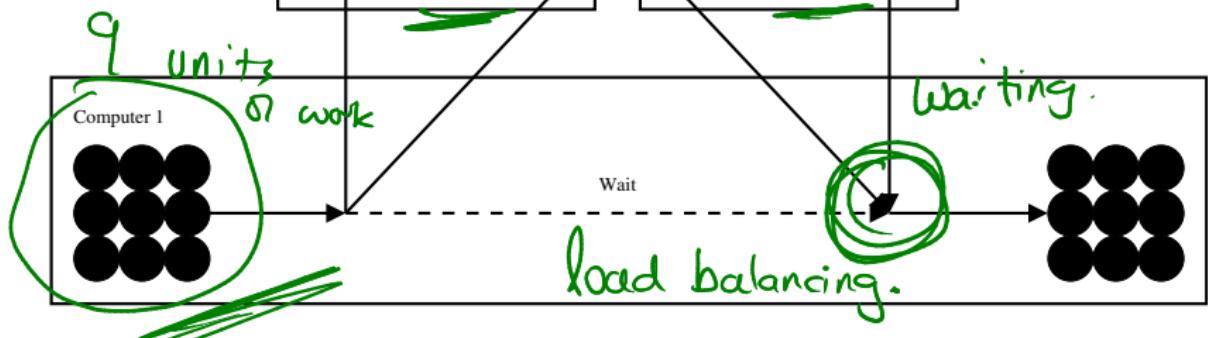
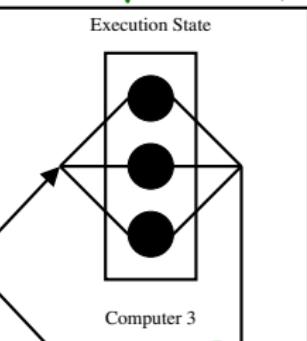


Figure 7: Model of execution on a distributed system (non-client/server). This shows a hybrid. But usually developers don't like these setups.

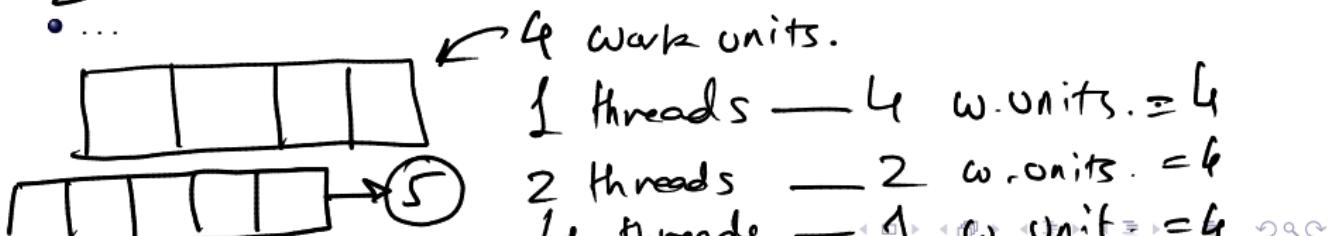
What is there in Parallel, that is not available in Serial?

Some Examples ...



total work → decompose into smaller tasks
 $\text{threads} \times \frac{\text{work done per thread}}$

- **Task Allocation:** Breakdown of total workload into smaller tasks, such that each smaller task is assigned to a different processor
- **Interim Communication:** Since each processor is working on independent tasks, they may need to communicate outcome of computation with other processors.
 - Is communication automatic? or defined by programmer?
 - Depends on type of multi-threading tool that is being used
- **Synchronization:** A waiting interval by a processor for certain tasks
 - E.g., arrival of data (before computation on it can be performed, or before bringing printed).
 - This would have positive effect on accuracy but negative effect on performance.
- ...



Applications and Examples

Task Parallel Threads

General Purpose Programs and Purposes

- Word processor (spell check, auto save, ...)
- Spreadsheet (automatic background recalculations after cell edits)
- Operating system (User threads, Kernel threads, ...)
- Server (multi client handling, database servers serving millions of transactions per second)
- GUI's (Event listeners)
- Video games (rendering of animation, taking care of physics, artificial intelligence, network connections, user I/O, ...))
- Web browser (tabs, multi-segment download accelerators, ...)
- Search Engines (Information searching, retrieval)
- Encoding/Decoding of video (compression/decompression)
- ...

threads are
doing different
types of
works.

Applications and Examples (cont.)

Scientific Purposes

→ Data Parallel
Threads -

- Mathematical problems (Ordinary differential equations, partial differential equations, linear algebra, numerical analysis)
 - Climate modelling and weather prediction
 - Fluid dynamics (video games, simulation of dam bursts, tsunamies, etc.)
 - Computational Biology, Finance, Physics, Chemistry
 - Machine learning, statistical methods
 - Image processing (medical images, spectral images, satellite images)
- Finding Aliens, Finding cure for Alzheimers, Cancers, Parkinson's, ...



Figure 8: The Folding@Home project by Stanford University

Data Parallel example

$$\begin{bmatrix} C_1 \\ C_2 \end{bmatrix} = \begin{bmatrix} A_1 \\ A_2 \end{bmatrix} + \begin{bmatrix} B_1 \\ B_2 \end{bmatrix}$$

for(i=0; i<2; i++) Removed

$$C[i] = A[i] + B[i]$$

i=Create two threads

$$C[i] = A[i] + B[i]$$

end two threads

multiple
threads

doing the
same task
on diff. locations

task parallel variation

$$C_1 = A_1 + B_1$$

$$C_2 = A_2 - B_2$$

for($i=0$; $i < 2$; $i++$)

$$\text{if}(i \% 2 == 0) C[i] = A[i] + B[i]$$

$$\text{else} \quad C[i] = A[i] - B[i]$$

$i =$ Create two threads

$$\text{if}(i \% 2 == 0) C[i] = A[i] + B[i]$$

$$\text{else} \quad C[i] = A[i] - B[i]$$

end two threads

multiple
threads
different
tasks on
different
locations.

task parallel variation

$$C_1 = A_1 + B_1$$

$$C_1 = A_1 - B_1$$

{ for(i=0; i<2; i++)

$$\text{if}(i/2 == 0) C = A + B$$

$$\text{else} \quad C = A - B$$

i=Create two threads

$$\text{if}(i/2 == 0) C = A + B$$

$$\text{else} \quad C = A - B$$

end two threads

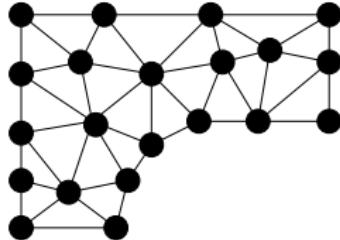
multiple
threads
different
tasks on
same location



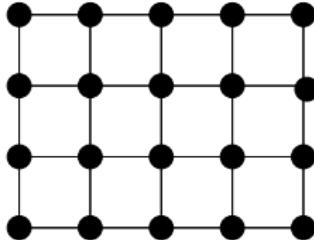
skip

Difference between both types of Applications

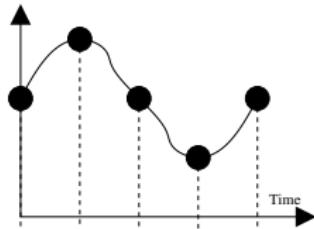
- General Purpose Applications: Threads are usually involved in doing **different things**. This is known as **task-parallelism**.
- Scientific Applications: Threads are usually involved in doing **similar things**, but at different locations. This is known as **data-parallelism**
- Different locations? Take a problem and perform both **spatial decomposition**, as well as **time decomposition**.
- Each decomposed location (involving predominantly local interactions) will be handled ideally by a single processor. If it involves non-local interactions, then processors would need to communicate.
- What level of decomposition would be required ?? A question answered by concerns of cost, speed, storage memory, functional requirements, etc.



Non-Uniform Spatial Decomposition



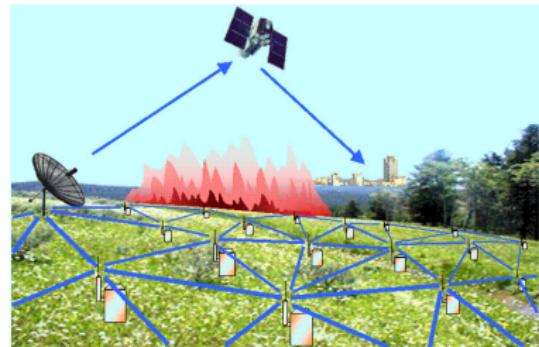
Uniform Spatial Decomposition



Uniform Time Decomposition

Applications

← Fig: Sensor Networks Deployment



- Approach-1: Client/Server Distributed Models
- Approach-2: Decentralized Distributed Models (P2P, WSN, MANET)
- Functions: Information Acquisition, Extraction, Aggregation and Control in geographically distributed systems
- E.g., sensor networks, where geographically distributed sensors **acquire** local information). This information is communicated to a control station through multiple hops. At each hop, information may be **extracted**, **aggregated** (to remove redundancies), and repackaged before forwarded. The control station ultimately receives the packet and performs **control** operations (e.g., fire control, smart cooling, etc.)
- Concerns: unreliable communication, device failures, cost, speed, deployment, energy constraints

Examples of Processors

Prog: task parallel vs

Data parallel.

- Hundreds of architectures, dozens in main-stream, hundreds now obsolete
- Many ways of classifying these architectures. The most simple classification is those that are “parallel”, and those that are “distributed”

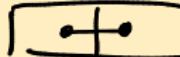
Array Processors (Vector Processors)



Stream
processors.

- Array of coupled processors (with nearest neighbor inter-connects)
- Most array processors perform “combined” or “group” execution.
- Ideal for spatial decomposed problems, mage processing, scientific computing, etc.
(not good for task-parallelism)

Data Parallel.



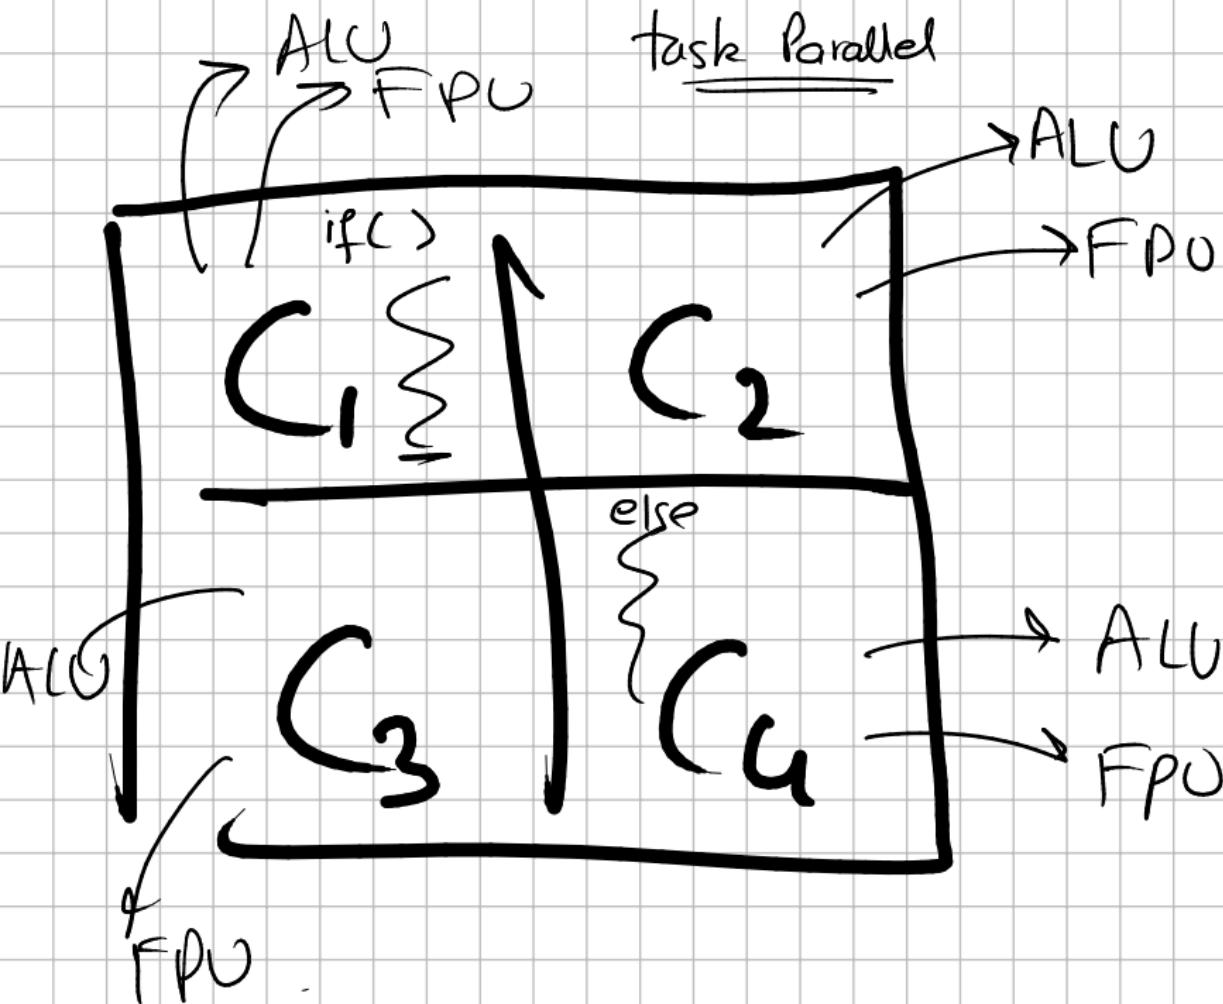
ALU
FPU

$$d_i = a_i + b_i$$

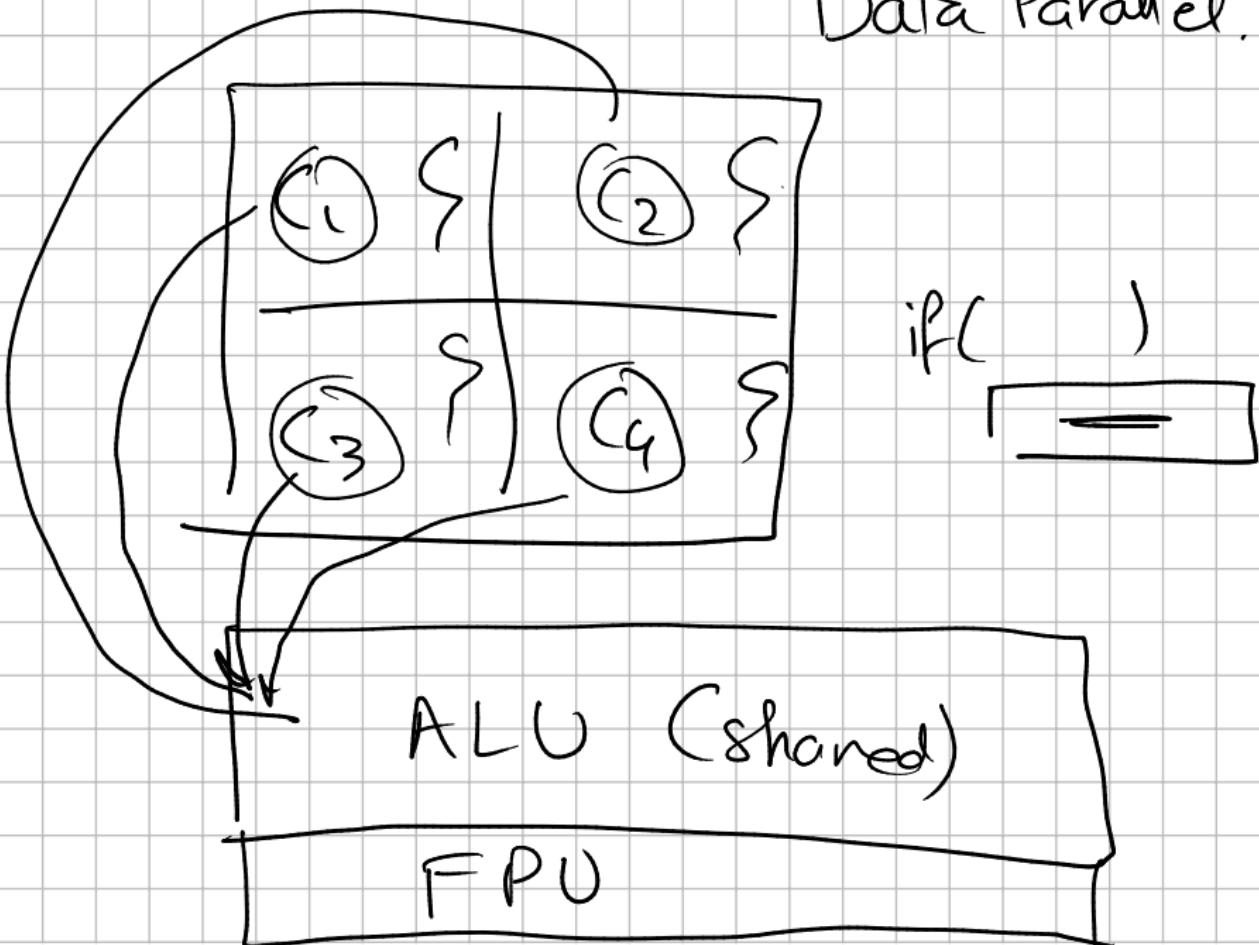
C ₁	C ₅
C ₂	C ₆
C ₃	C ₇
C ₄	C ₈

μP
8 cores.

array.

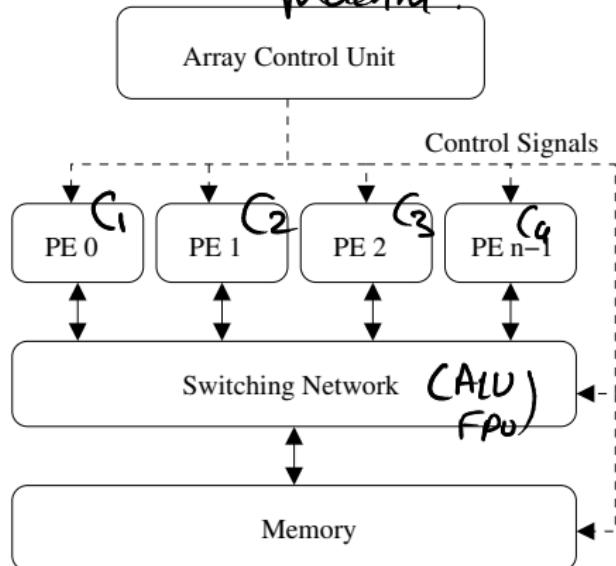


Data Parallel.



Examples of Processors (cont.)

Parallel Machine



Distributed Machine

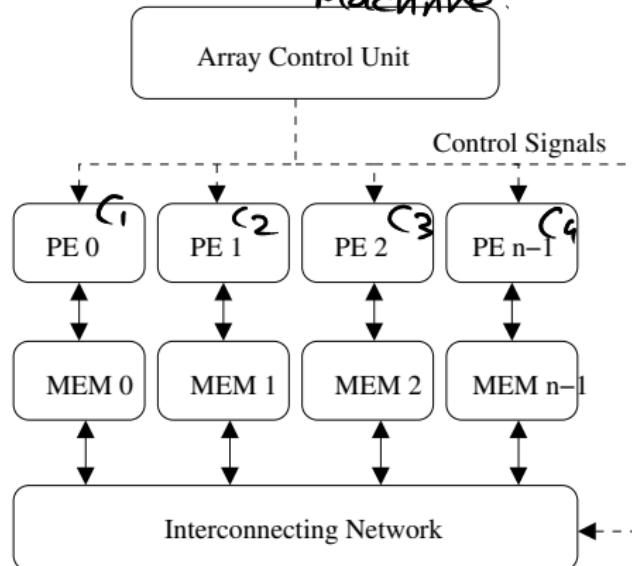


Figure 9: Array Processors (left) Shared Memory (right) distributed memory (IMG: Roland N Ibbett & Nigel P Topham, 1996) (Key: PE = Processing Element) \equiv Core.

Examples of Processors (cont.)

MP - multiprocessor.

Multiprocessors



Symmetric Multiprocessor.

- Loosely coupled processors, where each processor has more control on what to do and when to do.
- Well suited for task parallel problems (general purpose multi-threaded applications)
- Most common type is known as "symmetric multiprocessor system"

Systolic Arrays

→ Specialized Processor.

one + only one task. (ex. Fast)

- Makes use of "Very Large Scale Integration" (VLSI) Technology, where all computation elements are placed on a single chip. FPGA
- Processors are designed for a single purpose only and show fixed behaviour (code is embedded in system hardware)
- E.g., solving a linear system of equations, performing a fast Fourier transform
- Data "moves" across each processor (must be aligned before execution can commence)

Examples of Processors (cont.) Nof Important-

Matrix Multiplication.

$$\begin{array}{ccc}
 A_{00} & A_{01} & A_{02} & B_{00} & B_{01} & B_{02} \\
 A_{10} & A_{11} & A_{12} & x & B_{10} & B_{11} & B_{12} \\
 A_{20} & A_{21} & A_{22} & & B_{20} & B_{21} & B_{22}
 \end{array}
 = \boxed{A_{00}*B_{00} + A_{01}*B_{10} + A_{02}*B_{20}}$$

→ 3x3

$$\begin{array}{ccc}
 A_{00}*B_{01} + A_{01}*B_{11} + A_{02}*B_{21} & A_{00}*B_{02} + A_{01}*B_{12} + A_{02}*B_{22} \\
 A_{10}*B_{01} + A_{11}*B_{11} + A_{12}*B_{21} & A_{10}*B_{02} + A_{11}*B_{12} + A_{12}*B_{22} \\
 A_{20}*B_{01} + A_{21}*B_{11} + A_{22}*B_{21} & A_{20}*B_{02} + A_{21}*B_{12} + A_{22}*B_{22}
 \end{array}$$

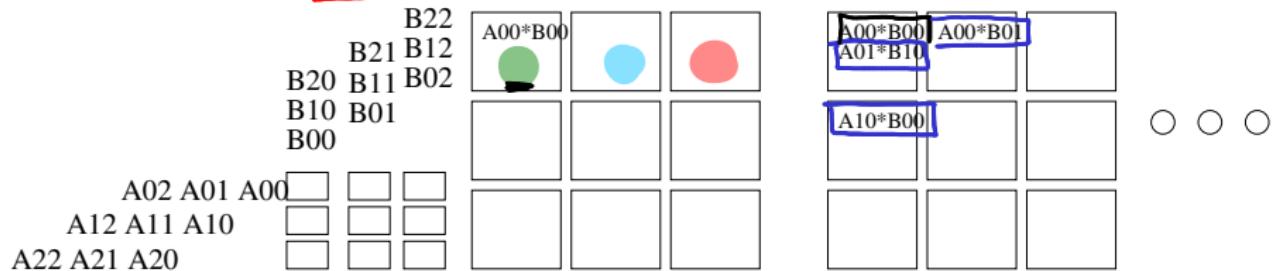
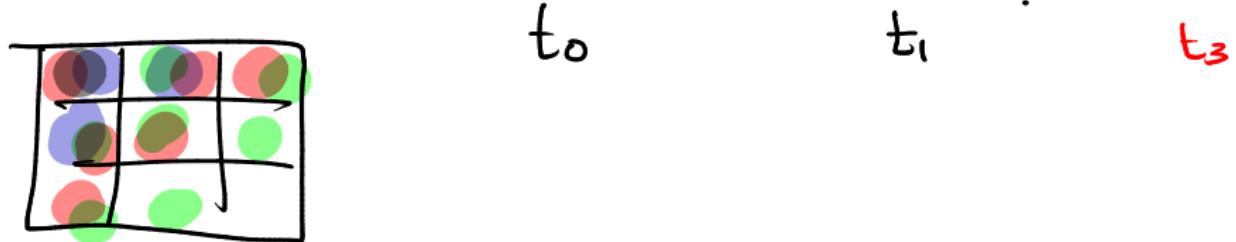


Figure 10: Movement of Data across systolic array for Matrix Multiplication



Flynn-Johnson Taxonomy (Classification)

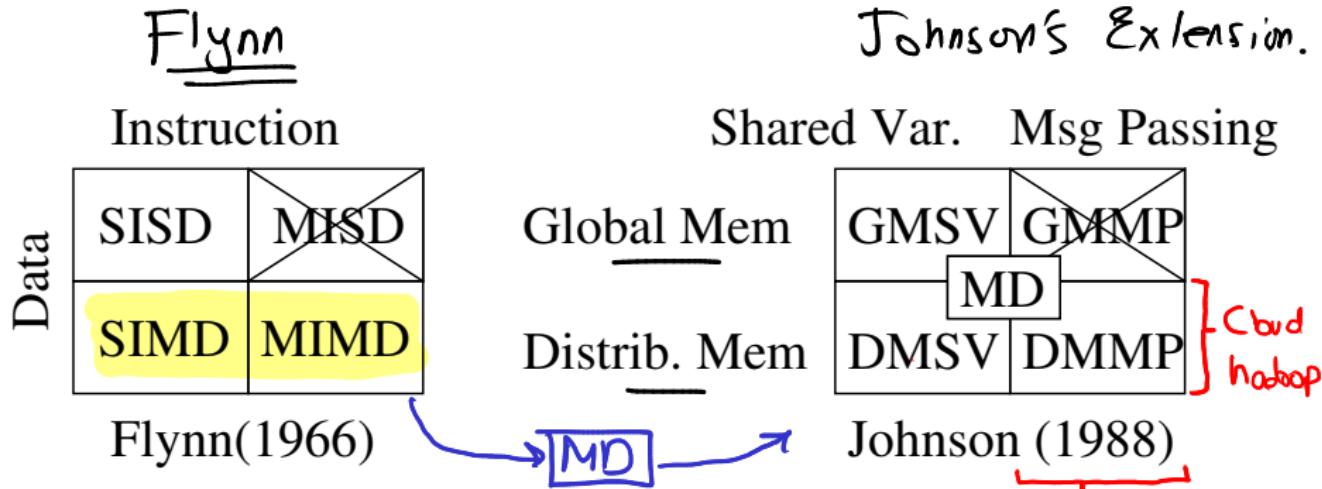


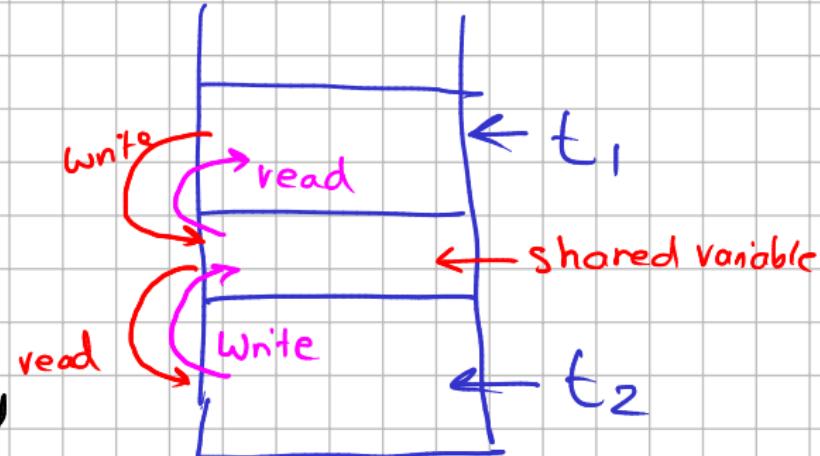
Figure 13: (l) Flynn's Classification (1966) based on instruction and data streams, and (r) Johnson's Classification (1988) based on memory and communication structures

- **Instruction Stream:** Sequence of Instructions
- **Data Stream:** Sequence of Data

Shared Variable

Synchronization

→ MPI → Distributed

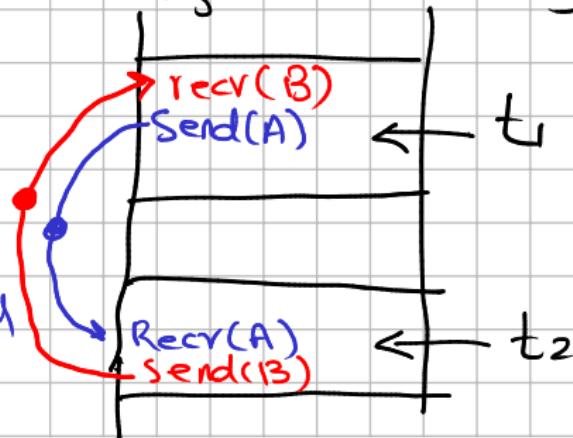


Message Passing → Blocking vs non-blocking

Send(t_1, t_2, A)

send(t_2, t_1, A)

↑
src dest.
↑
payload



Flynn

Code & Data

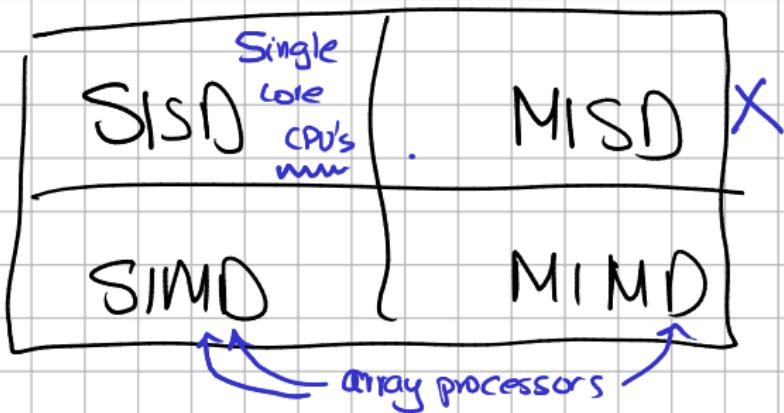
Code.
Instructions

SI

MI

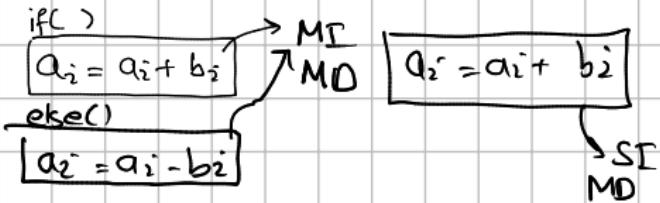
Data
Stream
—
MD

SD



Data.
A

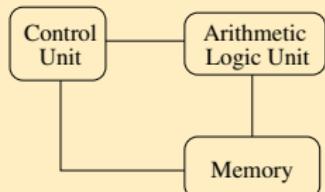
B



Flynn-Johnson Taxonomy (Classification) (cont.)

SISD

- Uni-Processor architectures
- Logic of program is sequential (sequential flow of instructions)
- Concurrent execution using context-switching can be possible
- *It's still SISD if you follow sequential programming on multi-core machines*



Typical Fetch/Execute Cycle: (1) Instruction Fetch, (2) Ins. Decode, (3) Operand Address Calculation, (4) Operand Fetch, (5) Execute, (6) Store Result

MISD

- Not commercially implemented
- E.g. multiple processors working on decrypting a single encrypted object (using different decrypting algorithms).

Flynn-Johnson Taxonomy (Classification) (cont.)

SIMD/MIMD

→ GPU's.

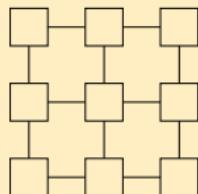
- Array/Streaming/Vector Processors interconnected using nearest-neighbour inter-connects.
- Processing elements contain only “Arithmetic Unit” instead of “Arithmetic and Logic Unit” (hence combined execution)
- Possible to disable some processing elements, if not needed
- Possible Issues
 - Data Alignment Issues
 - Conditional execution

GPU thread (deferring)

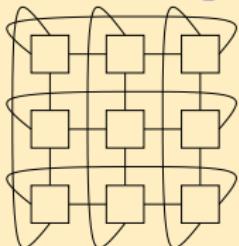
torus

hypercube

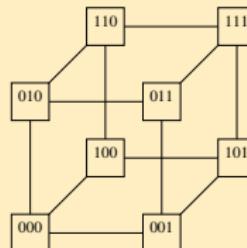
grid



Simple Nearest Neighbor
Inter-connects



3 by 3 Nearest Neighbor
Inter-connects



3-cube Inter-connects

MPI

Flynn-Johnson Taxonomy (Classification)

Johnson's Taxonomy

- **GMSV** Symmetric Multi-Processors (OpenMP, PThreads)
- **DMMP** Distributed systems (If simulated: Omnet++, NS2/3, etc.), NUMA based Cluster systems (OpenMPI, RPC), GPU Computing (OpenCL, CUDA), HPC based systems
- **DMSV** Cloud Model: Memory is distributed but access is through same address space (Hadoop, MapReduce)

(MD)

↳ Separate programming models.



1 Introduction

- Brief Overview
- Hardware Models

2 Shared Memory Programming Models: pThreads, OpenMP

- Overview
- Posix Threads
- First Look at OpenMP
- Profiling
- Work-Sharing Constructs
- Synchronization

3 Vector Programming: OpenCL/CUDA

- Overview
- GPGPU's: OpenCL & CUDA
- OpenCL Specification
- First Look at OpenCL

- CUDA Programming

- Optimization

4 Distributed Memory Programming Model: MPI

- Overview
- Communicators
- First Look at OpenMPI
- Sending/Receiving Messages
- Point-to-Point Send/Receive
- Collective Communication
- Multiple Communicators

5 Hadoop

- HDFS
- HDFS API's
- Map Reduce Framework

6 Spark