


# Work Sharing

- **Work Sharing:** General term describing distribution of work across threads
- Can be performed using three constructs:
  - for construct (for data parallelism)
  - sections construct (for task parallelism)
  - tasks construct (for irregular problems, e.g. unbounded loops, recursive codes)

Programming models

SSSS



Data Parallel


for threads. SIMD

SSSS



task parallel

# Work Sharing



- **Work Sharing:** General term describing distribution of work across threads
- Can be performed using three constructs:
  - **for** construct (for data parallelism)
  - **sections** construct (for task parallelism)
  - **tasks** construct (for irregular problems, e.g. unbounded loops, recursive codes)

tricky

# Data Parallelism

..

- Assuming that there is data independence across loop iterations, try:

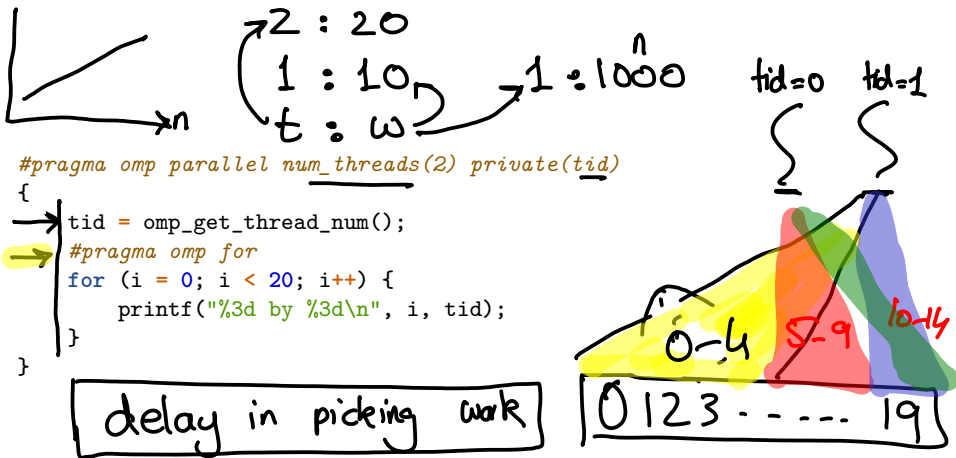
```
#pragma omp parallel [clauses]
{
    #pragma omp for [for clauses]
    for (loop control) {
        // statements
    }
}
```

→ Scheduling.

- OpenMP (or it's compilers) cannot (always) automatically identify data dependencies
- Threads “share” the iterations of the for loop
- Equivalent Code:

```
#pragma omp parallel for [for clauses]
for (loop control) {
    // statements
}
```

## Data Parallelism (cont.)



## Data Parallelism (cont.)

thread — workload.  
@ compile time

```
int chunksize = 5;
```

```
#pragma omp parallel num_threads(2) private(tid)
```

```
{
```

```
    tid = omp_get_thread_num();
```


```
    #pragma omp for schedule (static, chunksize)
```

```
    for (i = 0; i < 20; i++) {
```

```
        printf("%3d by %3d\n", i, tid);
```

```
    }
```

```
}
```



## Data Parallelism (cont.)

### Schedule Clauses (How loop iterations are mapped to threads)

#### Static Scheduling

- Low-overhead
- Load imbalance
- threads assigned “chunks” of iterations

#### Dynamic Scheduling


- High-overhead
- Reduces load imbalance
- Threads grab “chunks” of iterations

- `#pragma omp for schedule(static, chunksize)`
- `#pragma omp for schedule(dynamic, chunksize)`

$$6:24 = 1:4$$

Core=2

Static threads  $\leftarrow$   $\frac{\text{equally divided}}{\text{workload}}$




Dynamic threads  $\leftarrow$   $\frac{\text{Unequally divided}}{\text{workload}}$

(24)

1 : 1  
1 : 23

Core = 2

Static threads  $\leftarrow$  equally divided workload



as many.

---


Dynamic threads  $\leftarrow$  unequally divided workload



$$6 \div 24 = 1:6$$

Core=2

Static threads  $\leftarrow \frac{\text{equally divided}}{\text{workload}}$



Dynamic threads  $\leftarrow \frac{\text{Unequally divided}}{\text{workload}}$

Static threads  $\leftarrow \frac{\text{equally}}{\text{div'd}}$  workload

Dynamic threads  $\leftarrow \frac{\text{Unequally}}{\text{divided}}$  workload



## Data Parallelism (cont.)


• 1 — 1024 • 8 — 128  
 • 2 — 512 • 16 — 64 ..  
 • 4 — 256 • 32 — 32

•  $\boxed{64 - 16}$  • 512 2  
 • 128 — 8 • 1024 1  
 • 256 — 4



## Effect of Chunk Size and Thread Quantities


- Computed  $1024 \times 1024$  matrix multiplication using static and dynamic scheduling?
- (Left) Static Scheduling (Right) Dynamic Scheduling



Static

1024 threads  
1 chunks

better.




dynamic

threads  
chunksize

## Data Parallelism (cont.)

### Is Schedule Class Really Necessary?



**Fig. 1.** Speedup for different schedules on a 4-way IBM Power4

- Ayguade et. al., Is the schedule clause really necessary in OpenMP?, Technical Report, Springer Verlag, 2003

# Task Parallelism

- Considering following scenario:

```
p = pcibus();
n = networkCard(p);
w = wifiCard(p);
s = ssh(n,w);
h = http(n,w);
f = ftp(n,w);
```


- $n$ ,  $w$  can be executed in parallel
- $s$ ,  $h$ , and  $f$  can be executed in parallel
- Task Parallelism using sections in OpenMP:

```
#pragma omp parallel [clauses]
{
    #pragma omp sections
    {
        #pragma omp section
        {
            // Code of first task
        }
        #pragma omp section
        {
            // Code of second task
        }
    }
}
```

Num-threads (1)

← Sections

sharing  
th.  
work.



loop/I/O

## Task Parallelism (cont.)

- **Sections** must be inside a parallel region. **Sections** itself provides enclosure for an individual **omp section**

```
p = pcibus();
#pragma omp parallel sections num_threads(2)
{
    #pragma omp section
    n = networkCard(p);
    #pragma omp section
    w = wifiCard(p);
}
#pragma omp parallel sections num_threads(3)
{
    #pragma omp section
    s = ssh(n,w);
    #pragma omp section
    h = http(n,w);
    #pragma omp section
    f = ftp(n,w);
}
```

- If no. of threads is  $<$  no. of tasks, threads first attempt the beginning tasks before jumping to next ones.
- If no. of threads is  $>$  no. of tasks, each task is performed by one thread, while the remaining remain idle