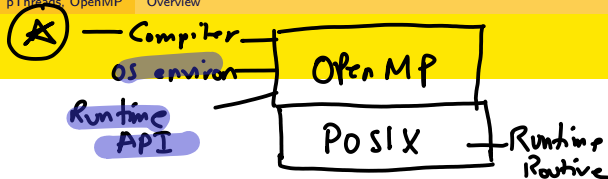


Overview



OpenMP

- Open Multi Processing
- Portable (Windows, Linux), shared-memory threading API for C/C++/Fortran (around 60 methods)
- Combines serial and parallel code in single source file
- Current specification: OpenMP 4.0

OpenMP

- Compiler directives (Native to C/C++/Fortran)
- Runtime library routines (e.g. increase/decrease threads as required)
- Environment variables

Compiler Directives

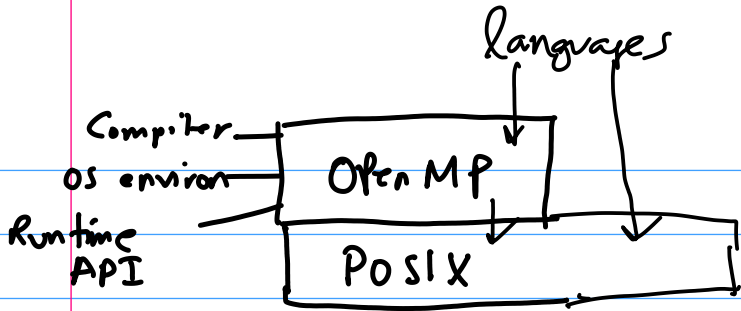
`#pragma omp construct clauses`

- `#pragma omp` Sentinel
- **construct** The construct, or directive Name
- **clauses** e.g., shared, schedule, etc.

native

```

#include
#define
#if
#ifndef
...
  
```



Overview (cont.)

Using

Runtime Routines

- 61 routines
- Perform tasks such as Control number of threads, Query thread information, lock management, wall clock time monitoring, etc.
- Example:

```
omp_set_num_threads(8);
```

Environment Variables

- Alternative to runtime routines

```
export OMP_NUM_THREADS=8 && source /etc/profile
```

X

X

OS

Compiler X

run time X

The Code (No. of Threads)

- Master thread spawns a team of threads as needed. (all thread creation work done transparently, developer saved from details)

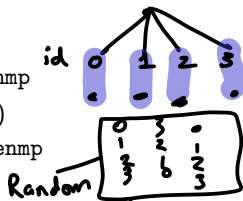
✱ `#include <omp.h>` // header

```
int main(int args, char *argv) {
    #pragma omp parallel num_threads(4)
    {
        tid = omp_get_thread_num();
        printf("Hello from %d\n", tid);
    }
}
```

✱ **Compiler** ✱

✱ **API** ✱

- Compilation (Intel)
`icc file.c -openmp`
- Compilation (GNU)
`gcc file.c -fopenmp`



- Common Runtime Routines:

```
omp_set_num_threads(int);
/* Who am I */
omp_get_thread_num();
omp_get_max_threads();
/* Are we in a parallel region? */
omp_in_parallel();
/* How many processing cores */
omp_get_num_procs();
/* Lock a thread */
omp_set_lock();
```

- Shared vs Private (What will happen below?)

```
int tid;
#pragma omp parallel private(tid)
{
    tid = omp_get_thread_num()
}
#pragma omp parallel shared(tid)
{
    tid = omp_get_thread_num()
}
```

int x → global.

```
void foo()
```

```
{
```

```
  int z
```

```
}
```

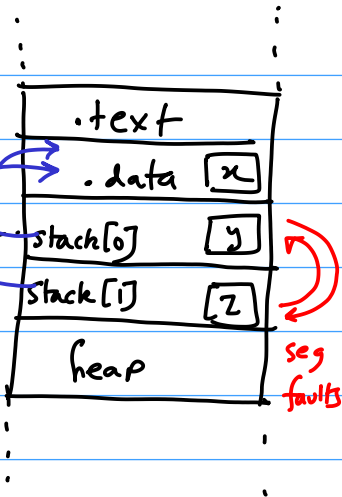
```
int main() ...
```

```
{
```

```
  int y
```

```
}
```

local



int x → global

int main()
{

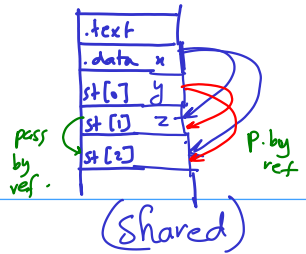
int y → local.

#pragma omp parallel shared(z)
{

int z

}

}



int x → global

```
int main()  
{
```

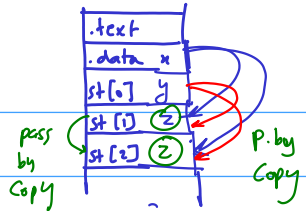
int y → local.

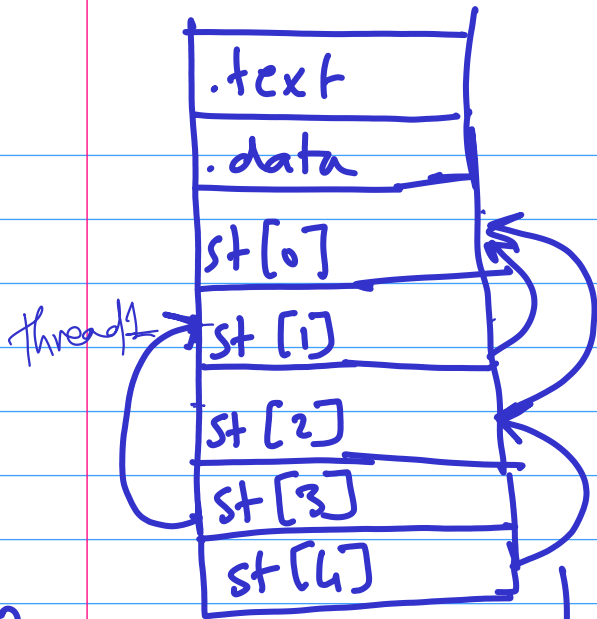
```
#pragma omp parallel private(z)  
{
```

int z

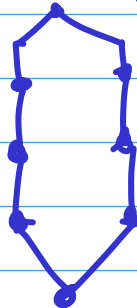
```
}
```

```
}
```





`foo()`
`main`



LIFO₂

LIFO₁

The Code (No. of Threads)

- Master thread spawns a team of threads as needed. (all thread creation work done transparently, developer saved from details)

```
#include <omp.h>
```

```
int main(int args, char *argv) {
    #pragma omp parallel num_threads(4)
    {
        tid = omp_get_thread_num();
        printf("Hello from %d\n", tid);
    }
}
```

- Compilation (Intel)

```
icc file.c -openmp
```

- Compilation (GNU)

```
gcc file.c -fopenmp
```

info
algorithms
parallel
enable.

- Common Runtime Routines:

```
omp_set_num_threads(int);
/* Who am I */
omp_get_thread_num();
omp_get_max_threads();
/* Are we in a parallel region? */
omp_in_parallel();
/* How many processing cores */
omp_get_num_procs();
/* Lock a thread */
omp_set_lock();
```

- Shared vs Private (What will happen below?)

```
int tid;
#pragma omp parallel private(tid)
{
    tid = omp_get_thread_num()
}
#pragma omp parallel shared(tid)
{
    tid = omp_get_thread_num()
}
```

The Code (No. of Threads)

- Master thread spawns a team of threads as needed. (all thread creation work done transparently, developer saved from details)

```
#include <omp.h>
```

```
int main(int args, char *argv) {
    #pragma omp parallel num_threads(4)
    {
        tid = omp_get_thread_num();
        printf("Hello from %d\n", tid);
    }
}
```

- Compilation (Intel)

```
icc file.c -openmp
```

- Compilation (GNU)

```
gcc file.c -fopenmp
```

- Common Runtime Routines:

```
omp_set_num_threads(int); // Who am I */
omp_get_thread_num();
omp_get_max_threads();
/* Are we in a parallel region? */
omp_in_parallel();
/* How many processing cores */
omp_get_num_procs();
/* Lock a thread */
omp_set_lock();
```

- Shared vs Private (What will happen below?)

```
int tid;
#pragma omp parallel private(tid)
{
    tid = omp_get_thread_num()
}
#pragma omp parallel shared(tid)
{
    tid = omp_get_thread_num()
}
```

Handwritten: #pragma omp parallel

The Code (No. of Threads)

- Master thread spawns a team of threads as needed. (all thread creation work done transparently, developer saved from details)

```
#include <omp.h>
```

```
int main(int args, char *argv) {
    #pragma omp parallel num_threads(4)
    {
        tid = omp_get_thread_num();
        printf("Hello from %d\n", tid);
    }
}
```

- Compilation (Intel)

```
icc file.c -openmp
```

- Compilation (GNU)

```
gcc file.c -fopenmp
```

- Common Runtime Routines:

```
omp_set_num_threads(int);
/* Who am I */
omp_get_thread_num(); —id
omp_get_max_threads(); —total
/* Are we in a parallel region?*/
omp_in_parallel();
/* How many processing cores*/
omp_get_num_procs();
/* Lock a thread */
omp_set_lock();
```

- Shared vs Private (What will happen below?)

```
int tid;
#pragma omp parallel private(tid)
{
    tid = omp_get_thread_num()
}
#pragma omp parallel shared(tid)
{
    tid = omp_get_thread_num()
}
```

The Code (No. of Threads)

- Master thread spawns a team of threads as needed. (all thread creation work done transparently, developer saved from details)

```
#include <omp.h>

int main(int args, char *argv) {
    #pragma omp parallel num_threads(4)
    {
        tid = omp_get_thread_num();
        printf("Hello from %d\n", tid);
    }
}
```

- Compilation (Intel)

icc file.c -openmp

- Compilation (GNU)

gcc file.c -fopenmp ←

- Common Runtime Routines:

```
omp_set_num_threads(int);
/* Who am I */
omp_get_thread_num();
omp_get_max_threads();
/* Are we in a parallel region? */
omp_in_parallel();
/* How many processing cores */
omp_get_num_procs();
/* Lock a thread */
omp_set_lock();
```

- Shared vs Private (What will happen below?)

```
int tid;
#pragma omp parallel private(tid)
{
    tid = omp_get_thread_num()
}
#pragma omp parallel shared(tid)
{
    tid = omp_get_thread_num()
}
```

0 dependence.

identification ✓

iterations ✓

```
for (i = 0; i < 5; i++)
    printf("hello %d", i);
```

hello 0

" 1

(?)

" 2

}

" 3

#pragma omp parallel num-threads(5)

int i = omp_get_thread_num();

3 printf("hello %d", i);

for loop \implies Omp parallel
Construct.

Scenario 1

~~for~~ (i)
 $A[i-1] + A[i-2]$

X

\rightarrow Special needs.

reduction \rightarrow reduce.

Scenario 1 = 2nd example

operators.

$+$ =

$-$ =

\times =

Sum = 0

for ($i = 0; i < 10; i++$)

Sum $\boxed{+ =}$ $a[i]$

operator
induced

$\times i-1, i-2$

previous

new.

Sum = $\textcircled{\text{Sum}} + a[i]$

Sum = 0
for (i = 0; i < 10; i++)
Sum += a[i]

- ① iterations ✓
- ② i index ✓

Sum = 0

option ①

```
#pragma omp parallel num_threads(10) shared(sum)
{
    int i = omp_get_thread_num();
    sum += a[i];
}
```

Nature of execution → Race Condition

```
load r1, a[i]
load r2, sum
add r1, r2, r3
store r3, sum
```


Sum = 0
for (i = 0; i < 10; i++)
Sum += a[i]


- ① iterations ✓
- ② i index ✓

Sum = 0

option ②

```
#pragma omp parallel num-threads(10) private(sum)
{
    int i = omp_get_thread_num()
    sum += a[i]
}
```

load r1, a[i]
load r2, sum
add r1, r2, r3
store r3, sum



Sum = 0

for($i=0; i < 10; i++$)

Sum += a[i]

- ① iterations ✓
- ② i index ✓

option ③

Sum = 0

#pragma omp parallel num-threads(10) reduction(sum: +)

{

int i = omp-get-thread-num()

Sum += a[i]

}

load r1, a[i]

load r2, sum

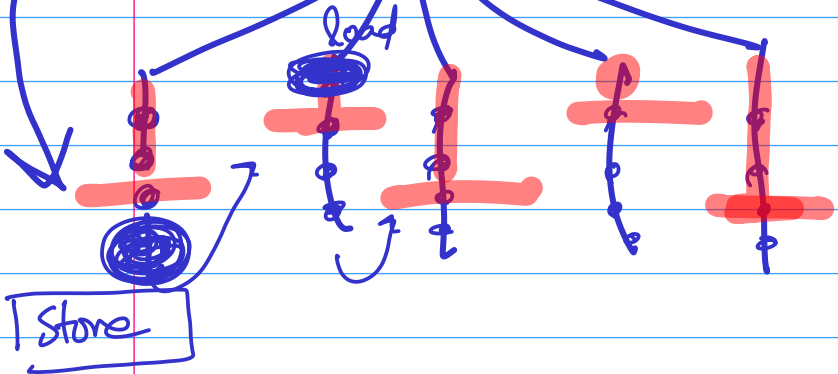
add r1, r2, r3

store r3, sum

pass by
reference

Shoved (sum)

Sum = 0



Pass copy

pirate

(sum)

Sum = 0

Sum

0

0

0

0

0

f =

all

0

↔

0

↔

0

↔

0

↔

0

option 3

reduction

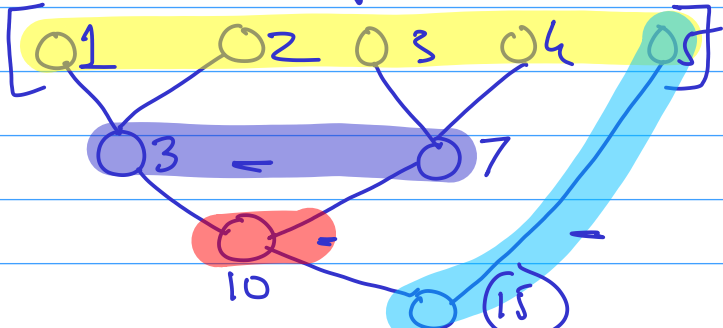
pass by copy

5 local copies

Sum = 0

option 2
+
option 1

5 threads
→
additional threads



for \Rightarrow threads
?

Scenario 2

while (1)
{
 recursive calls
 fork()
}
do/while

(?)

→

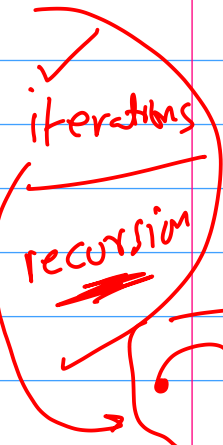
Scenario 3

fibonacci

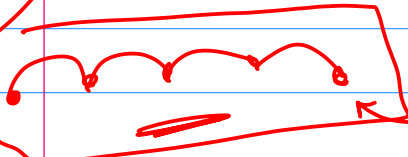
factorial (int a)

base condition

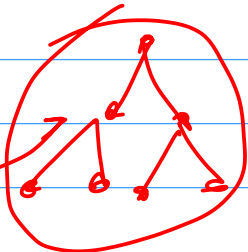
ret. factorial(a-1);



factorial (n);



thread



The Code (Reduction)

```
#include <omp.h>
#include <stdio.h>

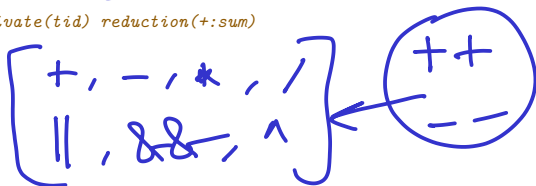
int main(int argc, char *argv)
{
    int tid, N = 1000, sum = 0;
    int array[N];

    for (i = 0; i < N; i++) {
        array[i] = i+1;
    }

    #pragma omp parallel num_threads(N) private(tid) reduction(+:sum)
    {
        tid = omp_get_thread_num();
        sum += array[tid];
    }

    printf("Sum: %d\n", sum);
}
```

reduction(~~Sum~~ : +)



The Code (Reduction)

```
#include <omp.h>
#include <stdio.h>
```

```
int main(int args, char *argv)
{
    int tid, N = 1000, sum = 0;
    int array[N];
```

```
    for (i = 0; i < N; i++) {
        array[i] = i+1;
    }
```

```
    #pragma omp parallel num_threads(N) private(tid) reduction(+:sum)
    {
        tid = omp_get_thread_num();
        sum += array[tid];
    }
```

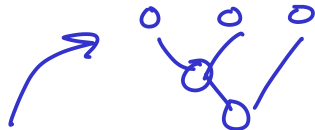
```
    printf("Sum: %d\n", sum);
}
```

algorithms v.v. fast

Handwritten diagram showing the reduction of two threads' values:

$$\begin{array}{c} 3 \\ \swarrow \text{true} \end{array} \parallel \begin{array}{c} 5 \\ \searrow \text{true} \end{array} = \textcircled{7}$$

$O(1)$



Handwritten diagram showing the reduction of two threads' values using a binary representation:

$$\begin{array}{r} 3 - 0011 \\ 5 - 0101 \\ \hline \boxed{0111} \end{array} \text{ OR}$$

The Code (Reduction)

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv)
{
    int tid, N = 1000, sum = 0;
    int array[N];

    for (i = 0; i < N; i++) {
        array[i] = i+1;
    }

    #pragma omp parallel num_threads(N) private(tid) reduction(+:sum)
    {
        tid = omp_get_thread_num();
        sum += array[tid];
    }

    printf("Sum: %d\n", sum);
}
```

load balancing.

tree structure
of threads

Complicated
example