




Syllabus

1 Distributed Memory Programming Model: MPI

- Overview
- Communicators



- First Look at OpenMPI
- Sending/Receiving Messages
- Point-to-Point Send/Receive
- Collective Communication
- Multiple Communicators





Single cpu ✓
multiple cpus (cluster) ✓

hidden

1 Distributed Memory Programming Model: MPI

- Overview
- Communicators

prog. model.



First Look at OpenMPI

- Sending/Receiving Messages
- Point-to-Point Send/Receive
- Collective Communication
- Multiple Communicators

message passing.

single
process

→ copies/duplicates

①

Open MPI



different
processes

②

Socket Programming




③


Remote Procedure Calls

single
process

//



RPC
Method Invocation



MPI Overview




- **Message Passing Interface**
- MPI Can be used for Shared Memory, as well as Distributed Memory architectures (Hybrid, if required)
- Supported by Fortran, C, C++ (but modules also available for python, & Java)
- Hides hardware details of underlying system (so portable)
- Many high performance libraries have MPI versions of API calls
- MPI version 3.0 specification has 400+ commands (function calls). Knowledge of only 11-12 of them can help you do the job in more than 90% of cases.

OpenMP.


Open MPI

MPI Communicators




MPI_COMM_WORLD: Name of default MPI Communicator

- A communication universe (communication domain, communication group) for a group of processes
- Stored in variables of type `MPI_COMM`
- Communicators are used as arguments to all message transfer MPI routines
- Each process within communicator has a rank; a unique integer identifier ranging between $[0, \#processors - 1]$
- Multiple communicators can be established in a single MPI program
- Intra-Communicator: Used for communication within a single group
- Inter-Communicator: Used for communication between two disjoint groups



global Communicator.

OpenMP $\rightarrow \begin{matrix} \{ \\ \{ \\ \{ \end{matrix}$ for each thread $\text{omp-get-thread-id}()$
0 2 1 (id)




MPI (id) \rightarrow Rank

for(i)
for(j)

j ↓


i →



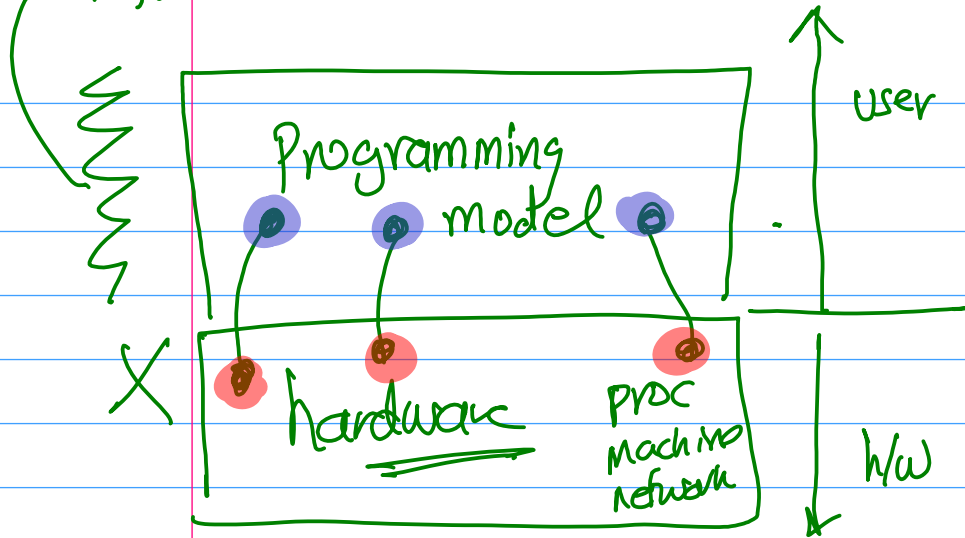
4 x 4

16 distributed
threads

MPI Communicators (cont.)




→ framework





MPI Communicators (cont.)



First Look (hellompi.c)

```
#include <stdio.h>
#include <mpi.h>
```

```
int main(int argc, char **argv)
{
```

```
    int size my_rank;
```

```
    ① MPI_Init(&argc, &argv);
```

```
    ② MPI_Comm_size(MPI_COMM_WORLD, &size); // check the total size
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
```

```
    printf("Hello from %d out of %d\n", my_rank, size);
```

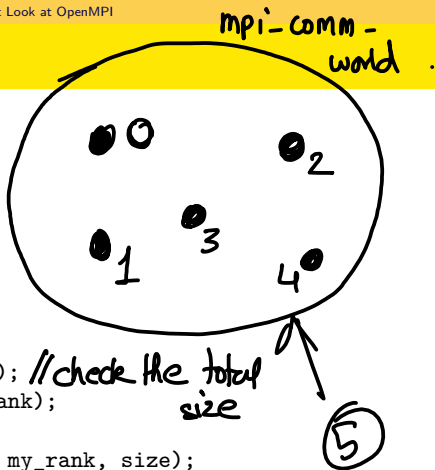
```
    MPI_Finalize();
```

```
    return 0;
```

```
}
```

```
mpicc hellompi.c # Compilation (mpicc for C++, also gcc hellompi.c -lmpi)
mpirun -np 4 -hostfile filename a.out # Execution
```

~~Set the total size~~



First Look (hellompi.c)

```
#include <stdio.h>
```

```
#include <mpi.h>
```

```
int main(int argc, char **argv)
```

```
{
```

```
    int size, my_rank;
```

```
    MPI_Init(&argc, &argv);
```

```
    MPI_Comm_size(MPI_COMM_WORLD, &size);
```

```
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
```

```
    printf("Hello from %d out of %d\n", my_rank, size);
```


```
    MPI_Finalize();
```

```
    return 0;
```


```
}
```

```
mpicc hellompi.c # Compilation (mpicc for C++, also gcc hellompi.c -lmpi)
```


```
mpirun -np 4 -hostfile filename a.out # Execution
```



initia



mpi-comm-world



First Look (hellompi.c)

```
#include <stdio.h>
```

```
#include <mpi.h>
```

```
int main(int argc, char **argv)
{
```

```
    int size, my_rank;
```

```
    MPI_Init(&argc, &argv);
```

```
    MPI_Comm_size(MPI_COMM_WORLD, &size);
```

```
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
```

```
    printf("Hello from %d out of %d\n", my_rank, size);
```


```
    MPI_Finalize();
```

```
    return 0;
```

```
}
```

```
mpicc hellompi.c # Compilation (mpicc for C++, also gcc hellompi.c -lmpi)
```


```
mpirun -np 4 -hostfile filename.a.out # Execution
```



ip1	1
ip2	3

Other activities.

C_2, C_3, C_4, b_3



First Look (hellompi.c)

```

#include <stdio.h>
#include <mpi.h>

int main(int argc, char **argv)
{
    int size, my_rank;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    printf("Hello from %d out of %d\n", my_rank, size);



    MPI_Finalize();

    return 0;
}

mpicc hellompi.c # Compilation (mpicc for C++, also gcc hellompi.c -lmpi)
mpirun -np 4 -hostfile filename a.out # Execution

```

Configuring a Simple MPI based Distributed Computing Cluster



Requirements

- SSH Server

`apt-get install openssh-server`

- OpenMPI Library

`apt-get install openmpi-bin openmpi-doc libopenmpi-dev`


- NFS Network File System

`apt-get install nfs-server nfs-client`


authentication SSH
password less

OpenMPI

mpich v2



/home/omar/temp/a.out



/home/omar/temp/a.out

Configuring a Simple MPI based Distributed Computing Cluster (cont.)



Passwordless Login

- Generate a public-private key pair for your user name. Do not specify a pass-phrase when asked

ssh-keygen ①

- Copy the public part of your key to the remote server

`ssh-copy-id <ip of remote computer>` ②

- If the above command does not work, continue with these commands:

- Copy the public key to the remote computer. For e.g., if the username is omar, the remote ip address is 1.2.3.4, then:

`scp /home/omar/.ssh/id_rsa.pub omar@1.2.3.4:/home/omar` Ⓐ

- Login to the remote computer using your username

`ssh 1.2.3.4 -l omar`

- Add the public key to authorized keys

`cat /home/omar/id_rsa.pub >> /home/omar/.ssh/authorized_keys`

Configuring a Simple MPI based Distributed Computing Cluster (cont.)

NFS

Dropbox

Manual copy

Transferring Files

- There are many ways to transfer files. You can setup an **NFS** mountpoint, share files using dropbox, or send files using scp. The scp method is given below:
`scp /location/of/a.out username@ipaddress:/home/username/a.out`
- Note:** All cluster nodes must be able to find the executable file at the same location as any other cluster node

